

Machine Learning with TensorFlow Project

Istanbul or Taipei?

Aslihan Demirkaya & Jerry Hsu



Project Objective:

This project is aimed to train the machine to differentiate between the cities "Istanbul" and "Taipei" when a picture of those cities is given as an input.

Ex: Input:

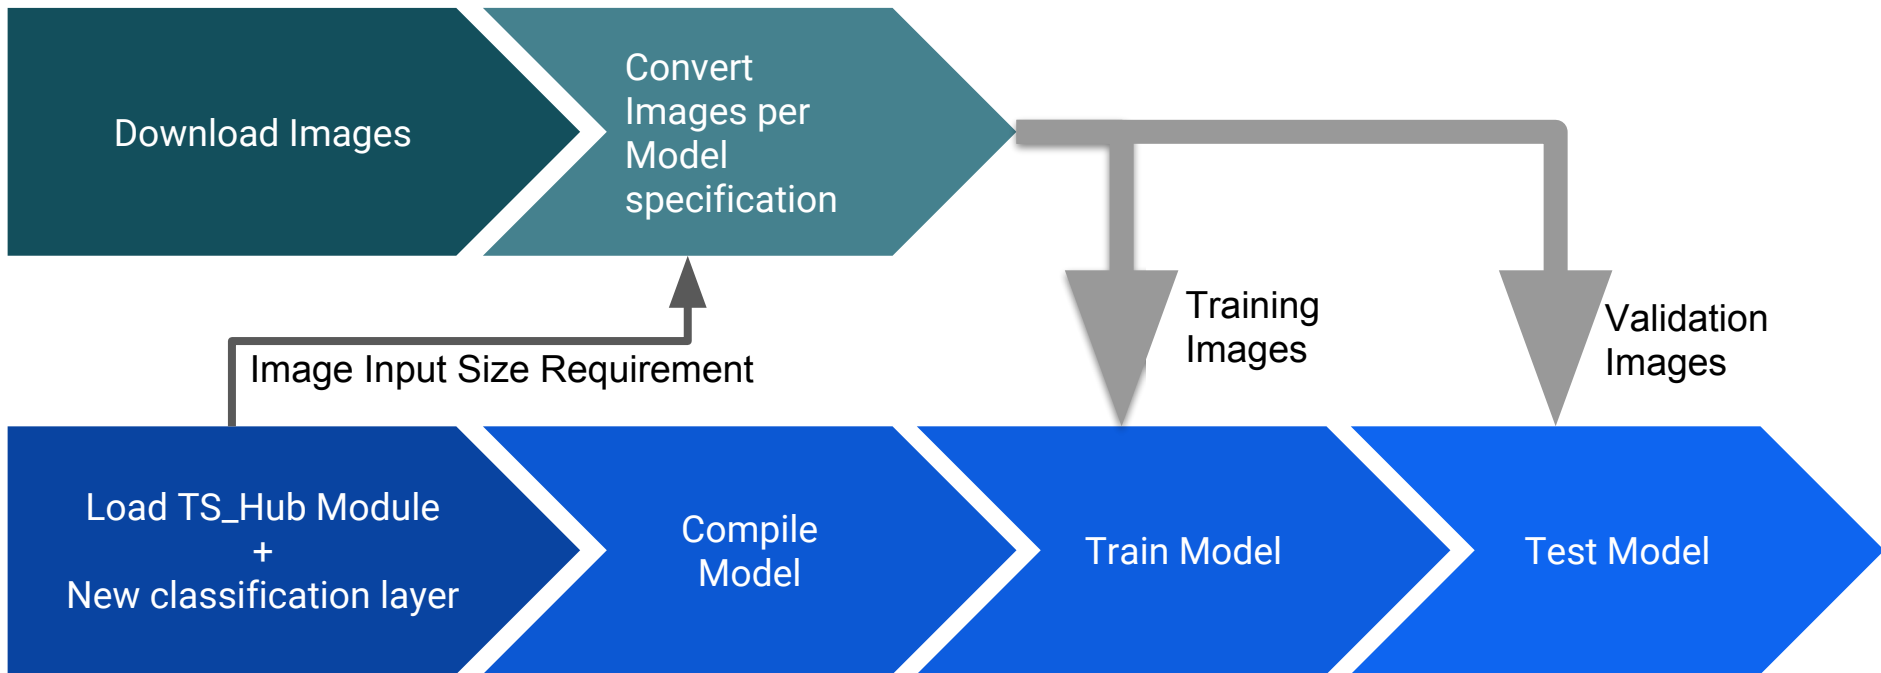


Can you guess the output?

Key Technology Components

- Use Flickr API and Images to acquire data
- Use Tensorflow Keras High Level API for rapid development
- Build network model by combining 'softmax' and MobileNetV2 module from Tensorflow Hub

Data Flow Overview



Download Images

Convert
Images per
Model Input
Spec

Load TS_Hub Module
+
New Classification layer

Compile
Model

Train Model

Test Model



Preparing the Data Set:

- We created an equally balanced data set consisting of images from both cities: Istanbul and Taipei. (3391 Istanbul + 3391 Taipei pictures)
- We used Flickr (an image hosting service) as the source of our data set.

- Steps to download the images:

Step 1: We install Python FlickrAPI from the Python Package Index:

<https://stuvel.eu/flickrapi-doc/>

Step 2: We created an app from this link:

<https://www.flickr.com/services/apps/create/noncommercial/> and obtained Flickr API key. The steps are explained in this video:

<https://www.youtube.com/watch?v=Lq1XR6dsDU>

Step 3: We download images of Istanbul and Taipei using the code below. The source for the code is found here:

<https://gist.github.com/yunjey/14e3a069ad2aa3adf72dee93a53117d6>

```
#In the first part, we obtain the urls of the images:
import flickrapi
import urllib
from PIL import Image

api_key = u'53a6e9e4eb5daab8bffbada47f0f4e4e8' # Flickr api access key is obtained in Step 2.
api_secret = u'dfb5411b99361e87'

flickr=flickrapi.FlickrAPI(api_key, api_secret, cache=True)

keyword = 'istanbul' #'taipei'

photos = flickr.walk(text=keyword,
                     tag_mode='all',
                     tags=keyword,
                     extras='url_c',
                     per_page=100, # Number of photos to return per page, the default number is 100.
                     sort='relevance')

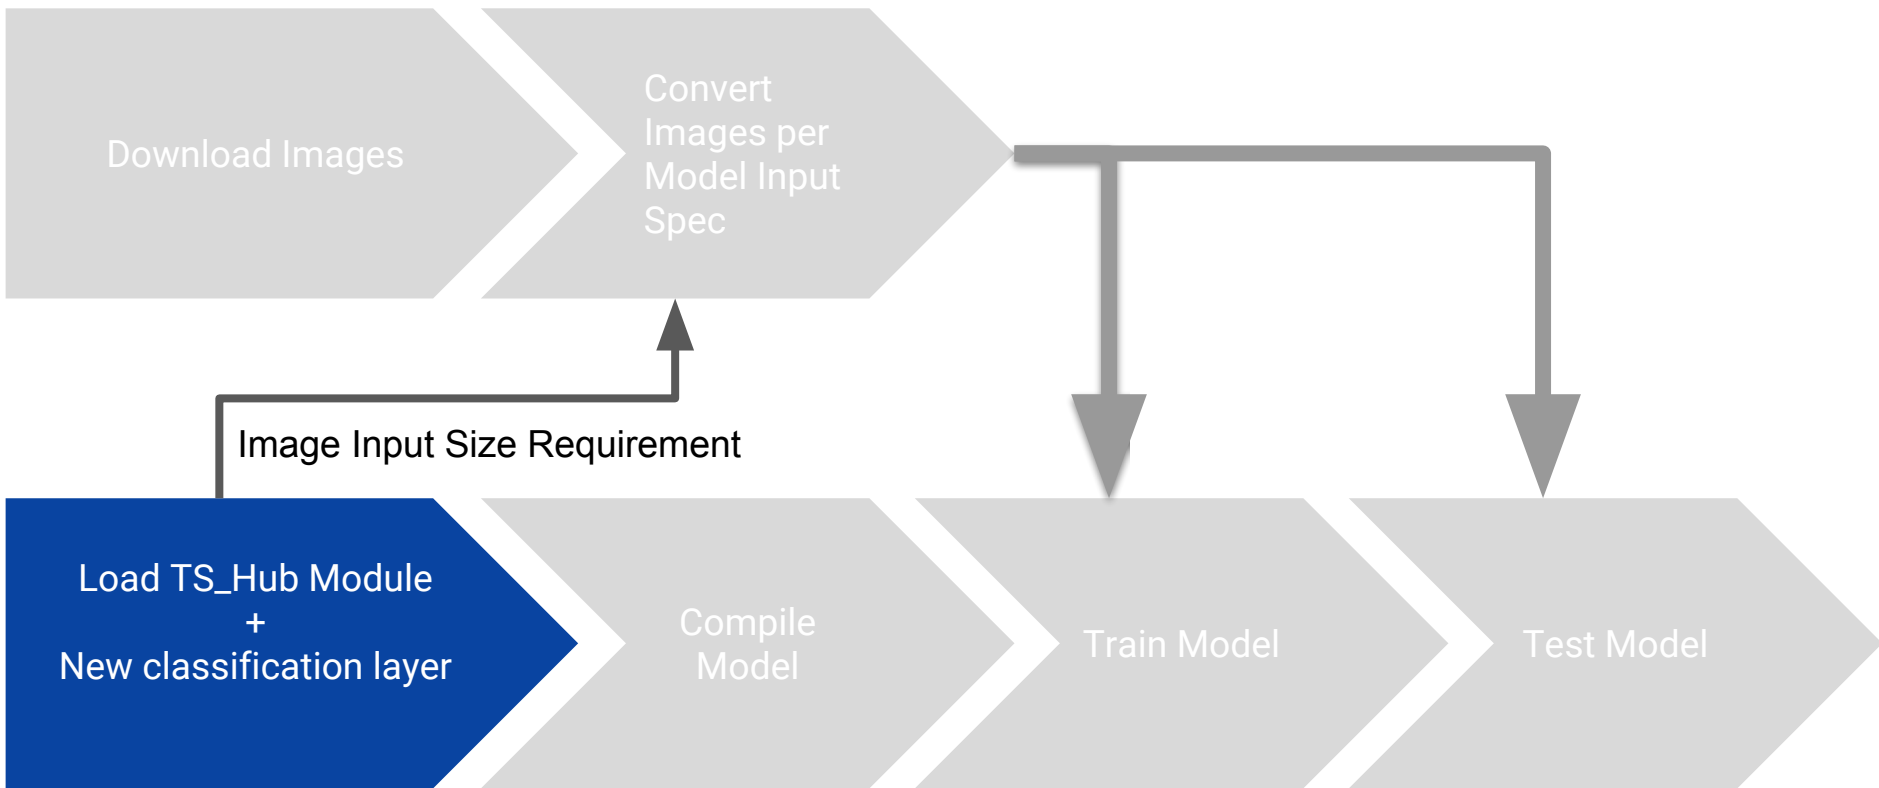
urls = []
for i, photo in enumerate(photos):

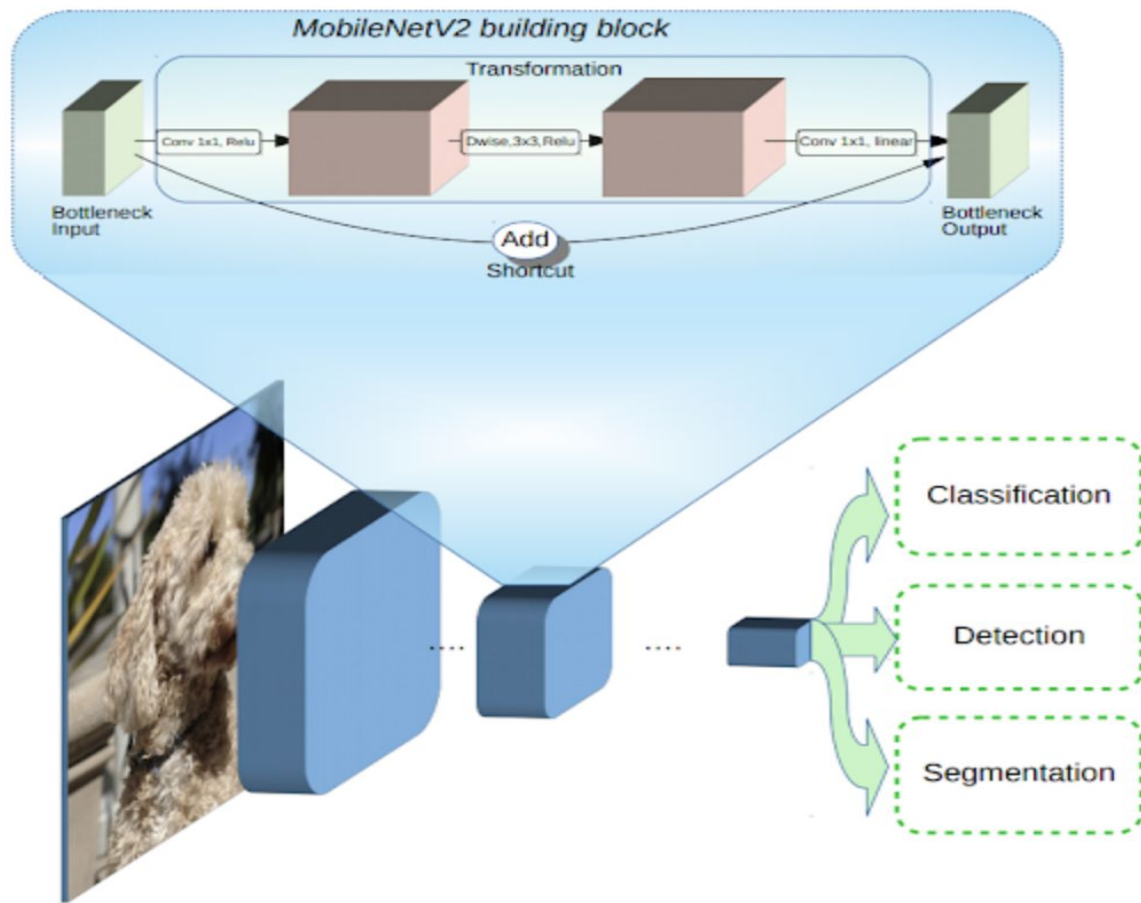
    url = photo.get('url_c')
    urls.append(url)

    if i > 4292: # get 4292 urls
        break
```


#In the second part we download images from the urls and resize them.

```
import pickle
pickle.dump(urls, open("istanbul_urls_flickr.pckl", 'wb'))
len(urls)
!mkdir istanbul
#pickle.dump(urls, open("taipei_urls_flickr.pckl", 'wb'))
#!mkdir taipei
for i, image_url in enumerate(urls):
    try:
        image = urllib.request.urlopen(image_url)
        # Resize the image and overwrite it
        image = Image.open(image)
        image = image.resize((256, 256), Image.ANTIALIAS)
        image.save(f'istanbul/{i}.jpg')
    except AttributeError:
        continue
```





<https://ai.googleblog.com/2018/04/mobilenetv2-next-generation-of-on.html>

Acquire training model and identify image size requirement

Using MobileNetV2 model without classification layer

```
feature_extractor_url = "https://tfhub.dev/google/imagenet/mobilenet_v2_100_224/feature_vector/2"

def feature_extractor(x):
    feature_extractor_module = hub.Module(feature_extractor_url)
    return feature_extractor_module(x)

IMAGE_SIZE = hub.get_expected_image_size(hub.Module(feature_extractor_url))
```

Prepare hub layer and create module

```
features_extractor_layer = layers.Lambda(feature_extractor, input_shape=IMAGE_SIZE+[3])
```

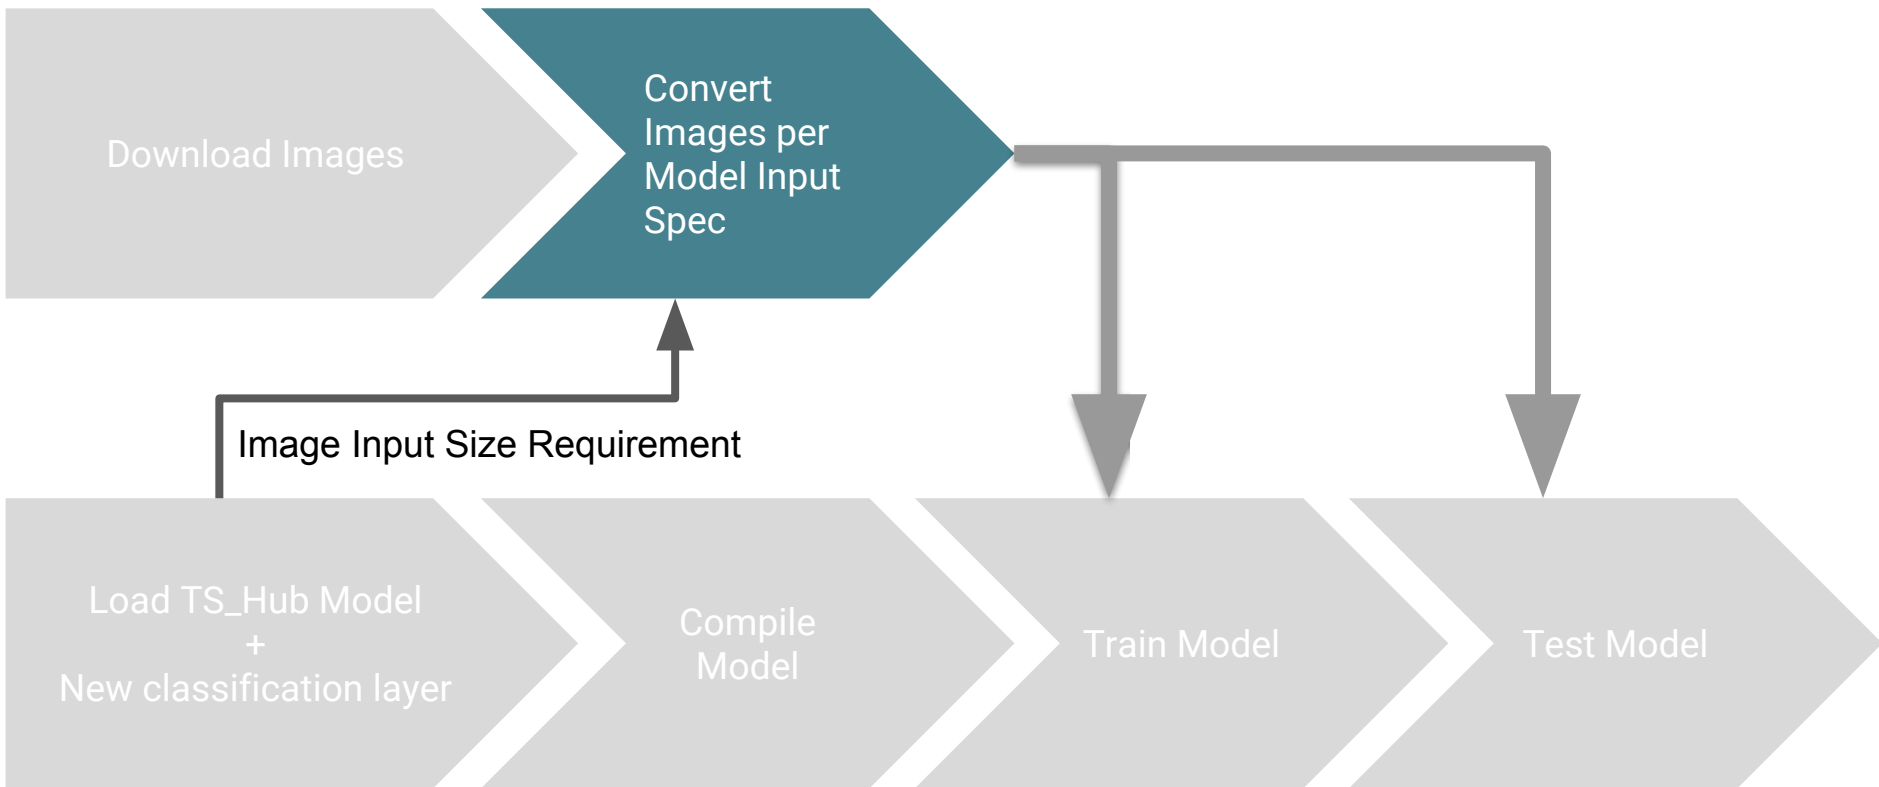
- Freeze the variables in the feature extractor layer, so that the training only modifies the new classifier layer.

```
features_extractor_layer.trainable = False
```

- Wrap the hub layer in a `tf.keras.Sequential` model, and add a new classification layer.

```
model = tf.keras.Sequential([
    features_extractor_layer,
    layers.Dense(training_data.num_classes, activation='softmax')
])
model.summary()
```

Layer (type)	Output Shape	Param #
=====		
lambda_1 (Lambda)	(None, 1280)	0 (~ 3M parameters)
dense_1 (Dense)	(None, 2)	2562
=====		
Total params: 2,562		
Trainable params: 2,562		
Non-trainable params: 0		



Setting up the data loader and splitting the data set into: **Training & Validation**

```
: data_root = 'istanbul_taipei'
  VALIDATION_SPLIT = 0.15
  image_generator = tf.keras.preprocessing.image.ImageDataGenerator(rescale=1/255, validation_split = VALIDATION_SPLIT)
  training_data = image_generator.flow_from_directory(str(data_root), target_size=IMAGE_SIZE, subset = 'training')

  validation_data = image_generator.flow_from_directory(str(data_root), target_size=IMAGE_SIZE
                                                       , subset = 'validation', shuffle = False)

  for image_batch, label_batch in training_data:
      print("Image batch shape: ", image_batch.shape)
      print("Label batch shape: ", label_batch.shape)
      break
```

Found 5766 images belonging to 2 classes.

Found 1016 images belonging to 2 classes.

Image batch shape: (32, 224, 224, 3)

Label batch shape: (32, 2)

Download Images

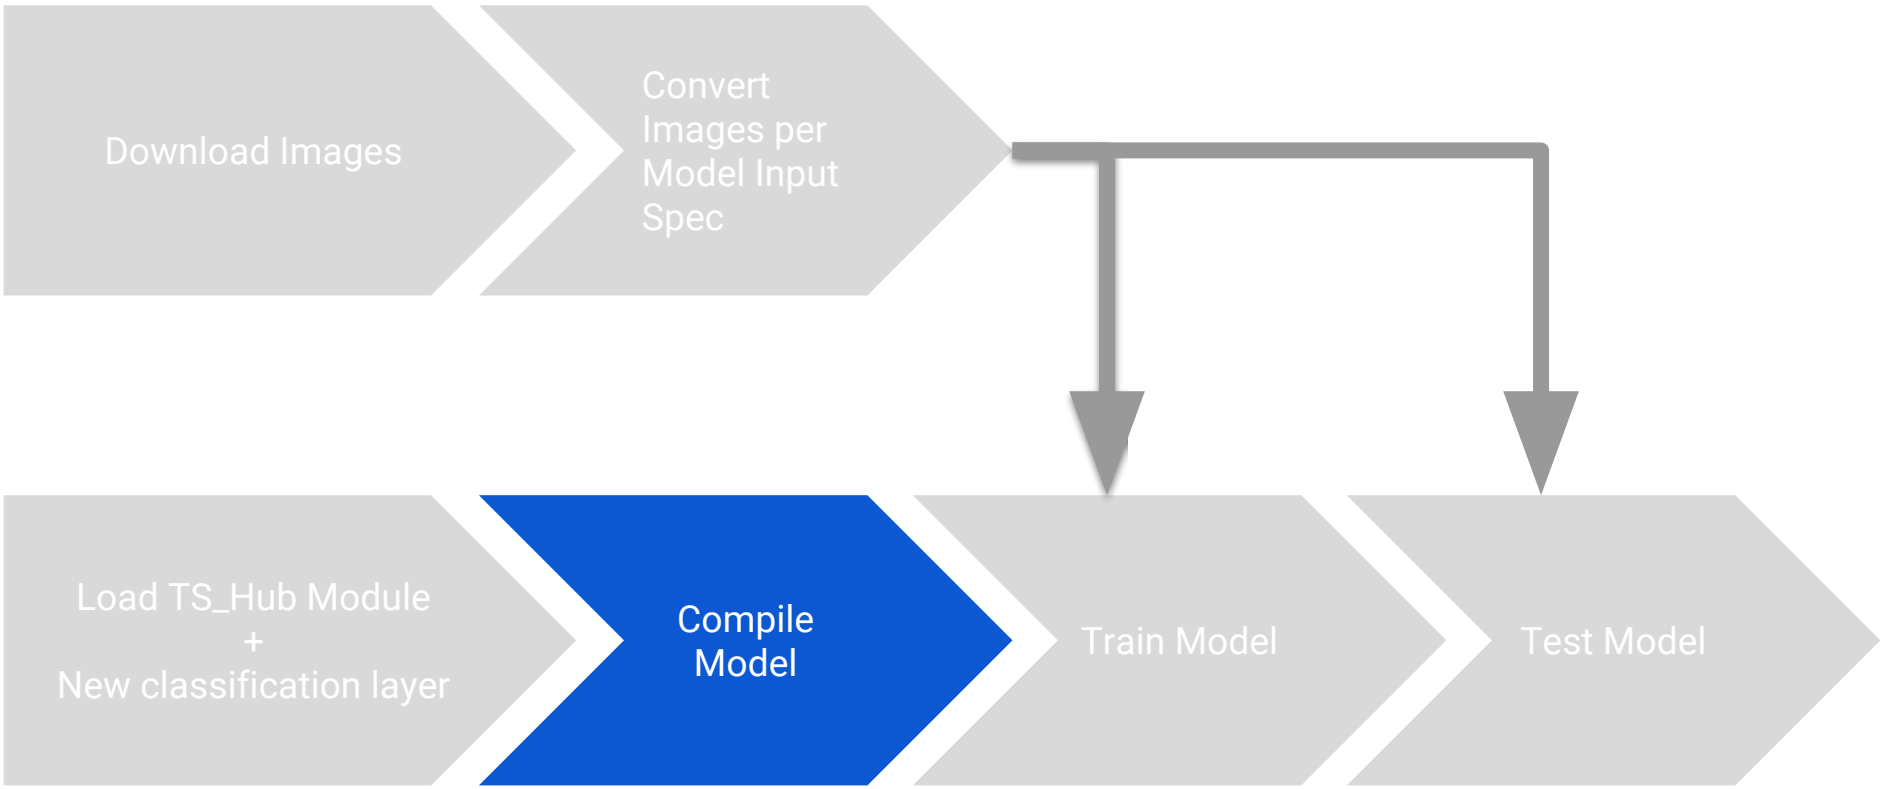
Convert
Images per
Model Input
Spec

Load TS_Hub Module
+
New classification layer

Compile
Model

Train Model

Test Model



Train the model

Use compile to configure the training process:

```
model.compile(  
    optimizer=tf.train.AdamOptimizer(),  
    loss='categorical_crossentropy',  
    metrics=[ 'accuracy' ] )
```

We use the fit method of Keras model to use custom Callback functions following:

https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/tutorials/images/hub_with_keras.ipynb#scrollTo=wC_AYRJU9NQe

```
class CollectBatchStats(tf.keras.callbacks.Callback):  
    def __init__(self):  
        self.batch_losses = []  
        self.batch_acc = []  
  
    def on_batch_end(self, batch, logs=None):  
        self.batch_losses.append(logs[ 'loss' ] )  
        self.batch_acc.append(logs[ 'acc' ] )
```

Download Images

Convert
Images per
Model Input
Spec

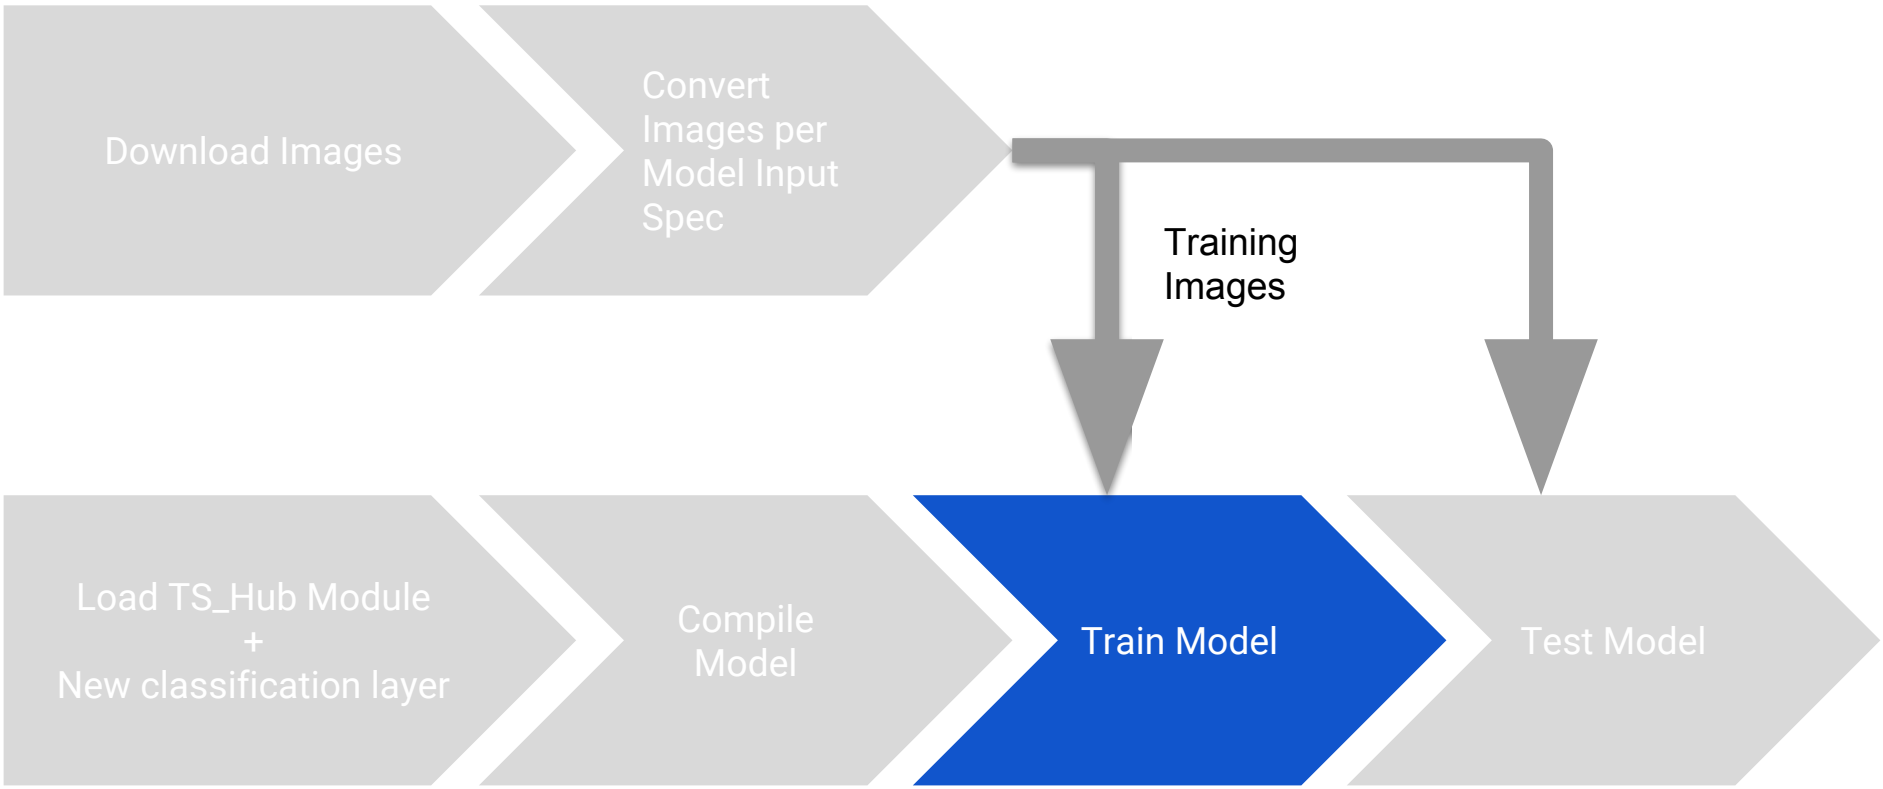
Training
Images

Load TS_Hub Module
+
New classification layer

Compile
Model

Train Model

Test Model



```
steps_per_epoch = training_data.samples//training_data.batch_size
batch_stats = CollectBatchStats()
model.fit((item for item in training_data), epochs=2, # Our experiments show that the accuracy does not improve after 2
          steps_per_epoch=steps_per_epoch,
          callbacks = [batch_stats])
```

Epoch 1/2

36/180 [=====>.....] - ETA: 1:34 - loss: 0.8016 - acc: 0.5877

Epoch 1/2

104/180 [=====>.....] - ETA: 52s - loss: 0.6407 - acc: 0.6722

Epoch 1/2

163/180 [======>...] - ETA: 11s - loss: 0.5866 - acc: 0.7013

Epoch 1/2

180/180 [=====] - 125s 692ms/step - loss: 0.5716 - acc: 0.7089

Epoch 2/2

21/180 [==>.....] - ETA: 1:46 - loss: 0.3903 - acc: 0.8229

Epoch 2/2

75/180 [=====>.....] - ETA: 1:12 - loss: 0.4139 - acc: 0.8025

Epoch 2/2


123/180 [=====>.....] - ETA: 39s - loss: 0.4192 - acc: 0.8039

Epoch 1/2


180/180 [=====] - 125s 692ms/step - loss: 0.5716 - acc: 0.7089

Epoch 2/2

180/180 [=====] - 124s 689ms/step - loss: 0.4139 - acc: 0.8075



```
plt.figure()  
plt.ylabel("Loss")  
plt.xlabel("Training Steps")  
plt.ylim([0,2])  
plt.plot(batch_stats.batch_losses)
```



```
plt.figure()  
plt.ylabel("Accuracy")  
plt.xlabel("Training Steps")  
plt.ylim([0,1])  
plt.plot(batch_stats.batch_acc)
```

Download Images

Convert
Images per
Model Input
Spec

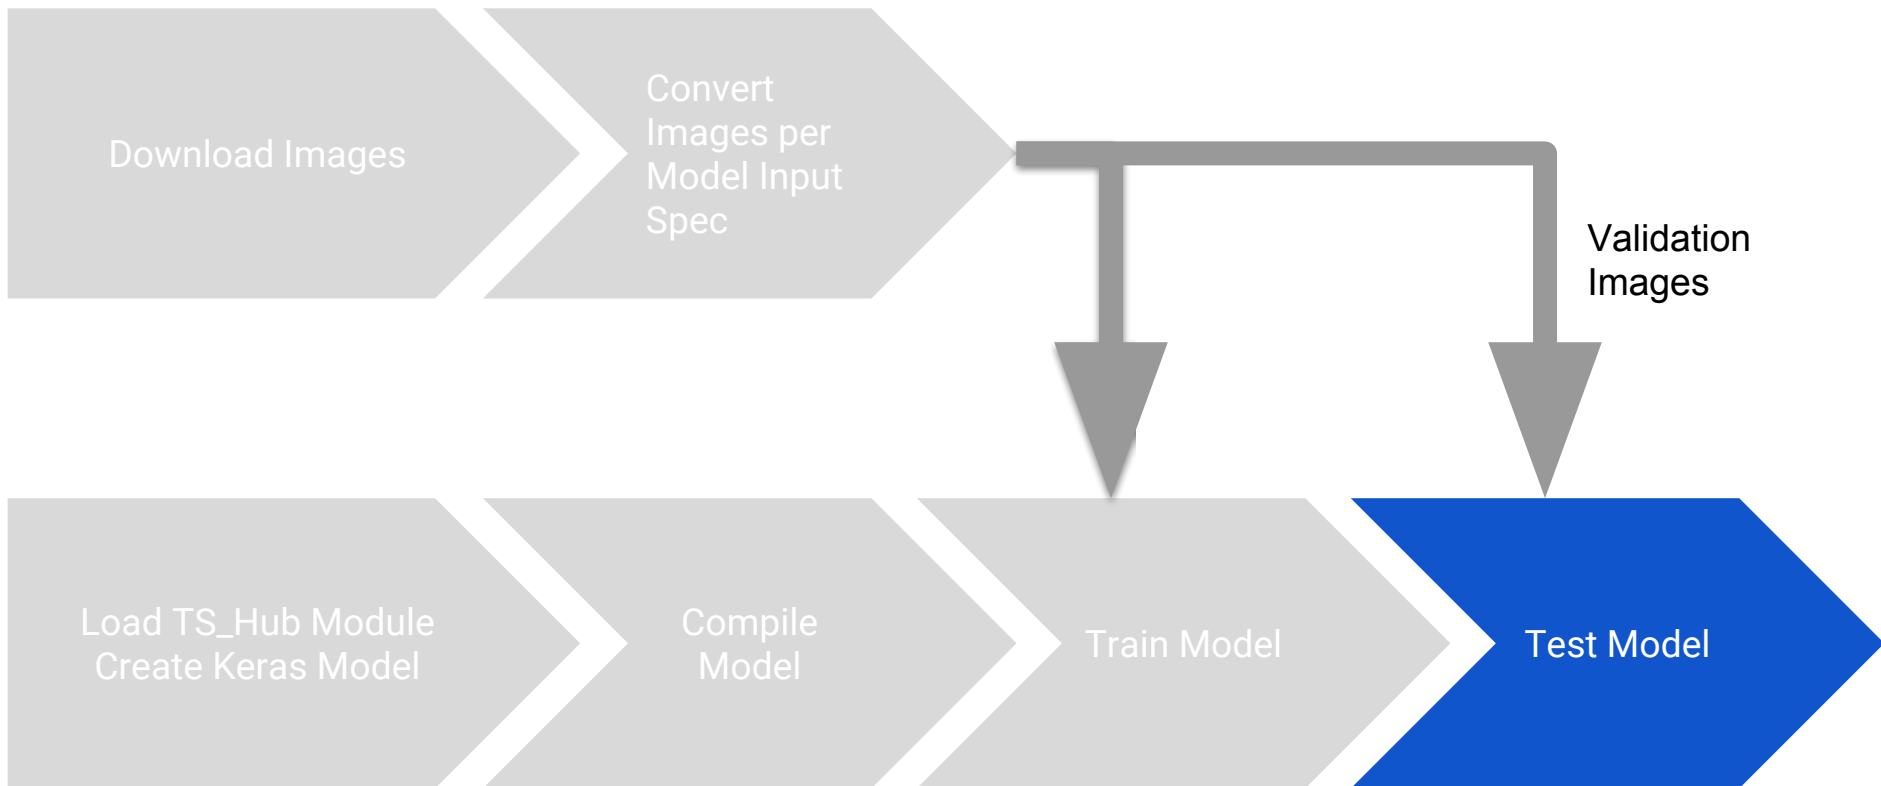
Validation
Images

Load TS_Hub Module
Create Keras Model

Compile
Model

Train Model

Test Model



Evaluating the model on the Validation Set

```
valid_scores = []
count_ = 0
num_batch = len(validation_data.filesnames) // 32
validation_data.reset() # resetting the generator so that we always get the same set of images for validation
for x_test, y_test in validation_data:
    scores = model.evaluate(x_test, y_test, verbose=1)
    valid_scores.append(scores)
    count_ += 1
    if count_ >= num_batch:
        break
print("validation accuracy", np.mean(list((uu[1] for uu in valid_scores))))
```


```
32/32 [=====] - 1s 21ms/sample - loss: 0.5485 - acc: 0.7500
32/32 [=====] - 1s 21ms/sample - loss: 0.7346 - acc: 0.6250
32/32 [=====] - 1s 21ms/sample - loss: 0.4502 - acc: 0.7812
32/32 [=====] - 1s 22ms/sample - loss: 0.4661 - acc: 0.8438
32/32 [=====] - 1s 22ms/sample - loss: 0.7179 - acc: 0.6562
32/32 [=====] - 1s 22ms/sample - loss: 0.3870 - acc: 0.7812
validation accuracy 0.8054435
```

Results:

pred: Istanbul0.91
actual: Istanbul



pred: Istanbul0.99
actual: Istanbul



pred: Istanbul0.96
actual: Istanbul



pred: Istanbul0.92
actual: Istanbul



pred: Istanbul0.66
actual: Istanbul



pred: Istanbul0.99
actual: Istanbul




pred: Istanbul0.89
actual: Istanbul



pred: Istanbul0.83
actual: Istanbul



pred: Istanbul0.99
actual: Istanbul



pred: Istanbul0.62
actual: Istanbul



pred: Istanbul0.96
actual: Istanbul



pred: Istanbul0.93
actual: Istanbul



pred: Istanbul0.96
actual: Istanbul



pred: Istanbul0.66
actual: Istanbul



pred: Istanbul0.96
actual: Istanbul




pred: Istanbul0.95
actual: Istanbul



pred: Taipei0.13
actual: Istanbul



pred: Istanbul0.98
actual: Istanbul



pred: Istanbul0.88
actual: Istanbul



pred: Istanbul0.72
actual: Istanbul



pred: Istanbul1.00
actual: Istanbul



pred: Istanbul0.78
actual: Istanbul



pred: Istanbul0.96
actual: Istanbul




pred: Istanbul0.99
actual: Istanbul



pred: Istanbul0.97
actual: Istanbul



pred: Taipei0.34
actual: Istanbul




pred: Istanbul0.89
actual: Istanbul



pred: Istanbul0.89
actual: Istanbul



pred: Taipei0.49
actual: Istanbul



pred: Istanbul0.74
actual: Istanbul



pred: Taipei0.21
actual: Taipei



pred: Taipei0.06
actual: Taipei



pred: Taipei0.06
actual: Taipei



pred: Taipei0.18
actual: Taipei



pred: Taipei0.15
actual: Taipei



pred: Istanbul0.51
actual: Taipei



pred: Taipei0.38
actual: Taipei



pred: Taipei0.17
actual: Taipei



pred: Taipei0.45
actual: Taipei



pred: Taipei0.21
actual: Taipei




pred: Taipei0.33
actual: Taipei



pred: Taipei0.06
actual: Taipei



pred: Istanbul0.64
actual: Taipei



pred: Taipei0.07
actual: Taipei



pred: Istanbul0.92
actual: Taipei



pred: Istanbul0.92
actual: Taipei



pred: Istanbul0.97
actual: Taipei



pred: Taipei0.05
actual: Taipei



pred: Taipei0.05
actual: Taipei



pred: Istanbul0.56
actual: Taipei



pred: Taipei0.11
actual: Taipei



pred: Taipei0.19
actual: Taipei



pred: Taipei0.43
actual: Taipei



pred: Taipei0.08
actual: Taipei



pred: Taipei0.28
actual: Taipei



pred: Istanbul0.63
actual: Taipei



pred: Taipei0.44
actual: Taipei



pred: Taipei0.50
actual: Taipei



pred: Istanbul0.78
actual: Taipei



pred: Taipei0.31
actual: Taipei



Confusion Matrix

```
predicted_istanbul=[None]*len(prob_istanbul)
actual_istanbul=[None]*len(true_istanbul)
for i in np.arange(len(true_istanbul)):
    if (prob_istanbul[i] >0.5): #threshold value is taken as 0.5
        predicted_istanbul[i]="Istanbul"
    else:
        predicted_istanbul[i]="Taipei"
for i in np.arange(len(true_istanbul)):
    if (true_istanbul[i]==1):
        actual_istanbul[i]="Istanbul"
    else:
        actual_istanbul[i]="Taipei"
```

```
import pandas as pd
actual = actual_istanbul
predicted = predicted_istanbul
df = pd.DataFrame({'predicted':predicted,'actual':actual})


pd.crosstab(df['predicted'],df['actual'])
```

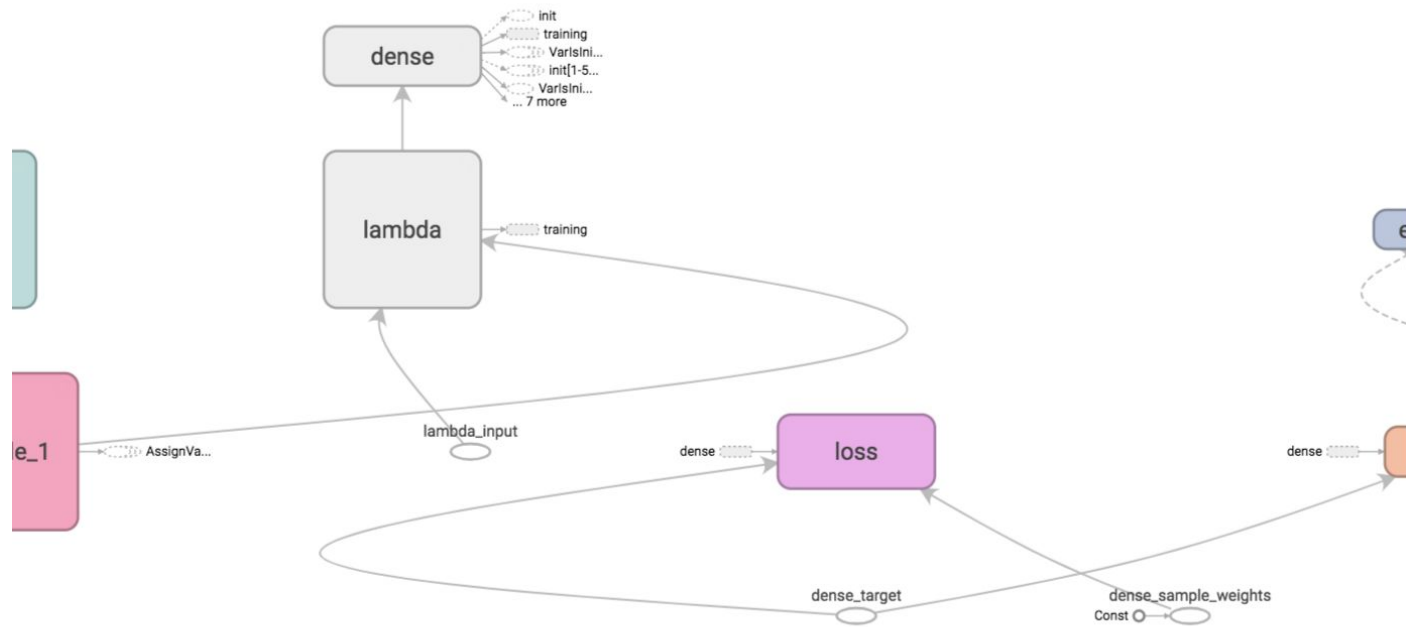
Confusion Matrix:

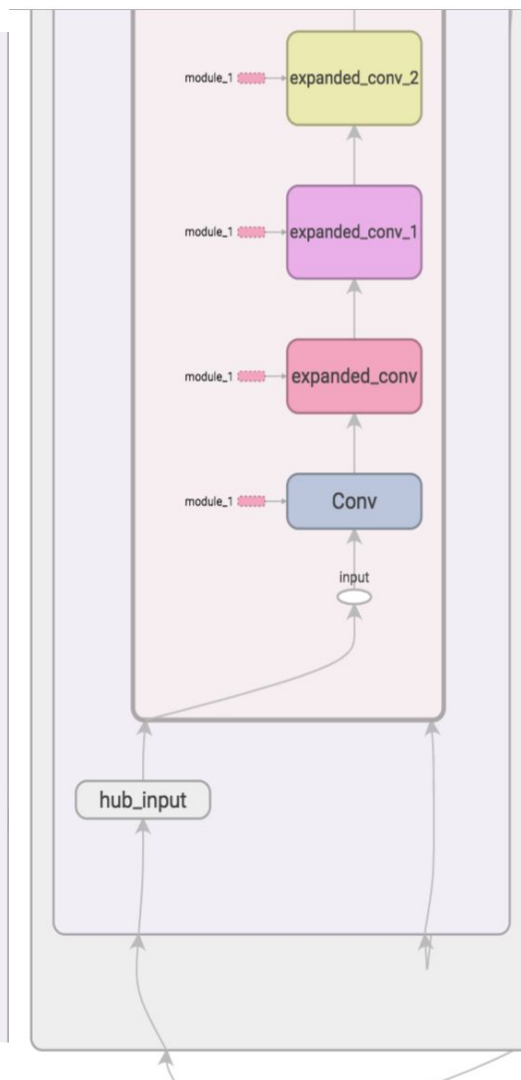
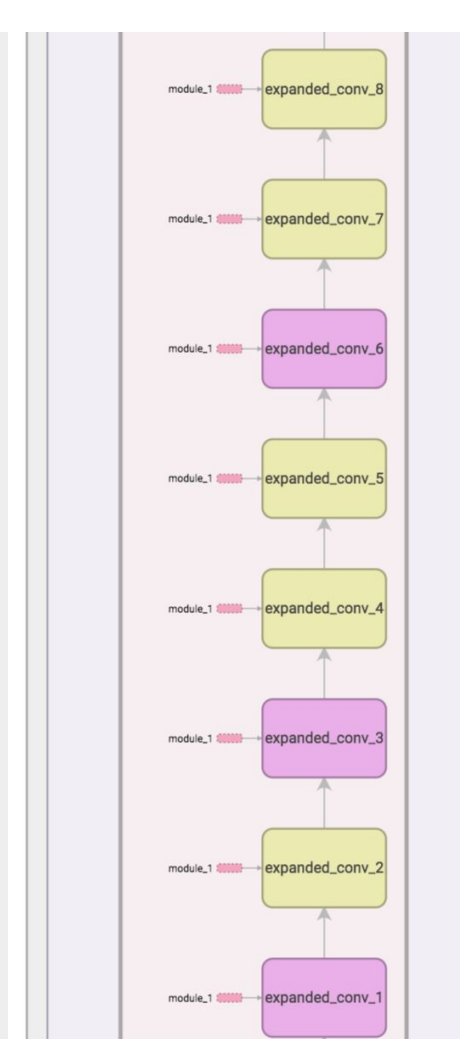
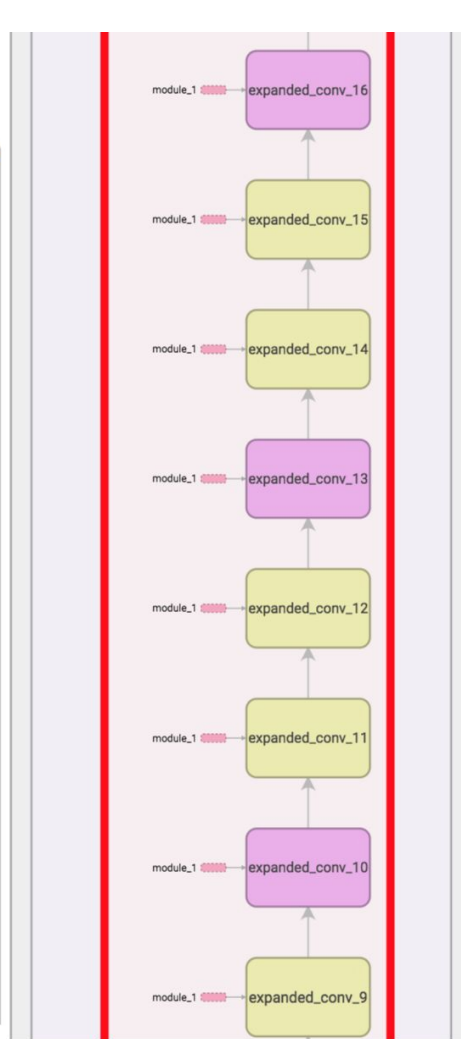
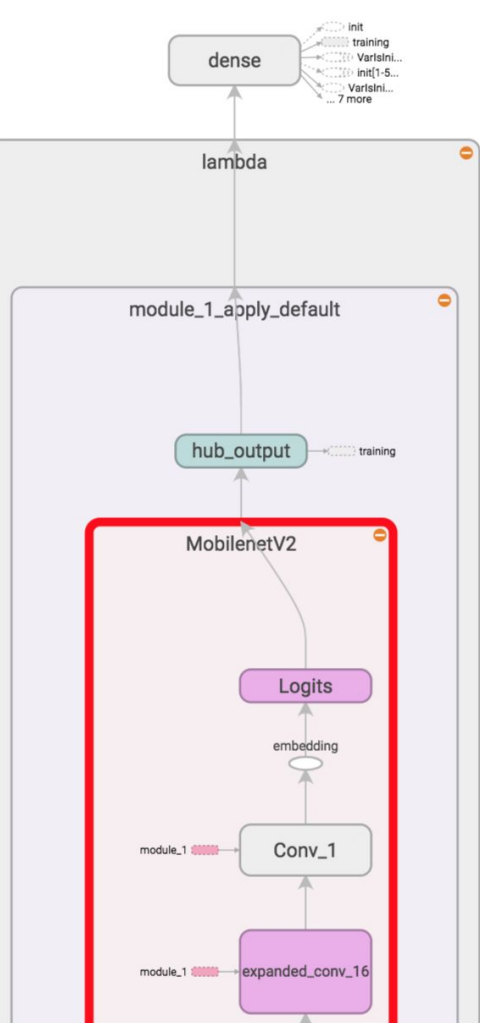
actual		Istanbul	Taipei
predicted	Istanbul	439	133
	Taipei	69	375

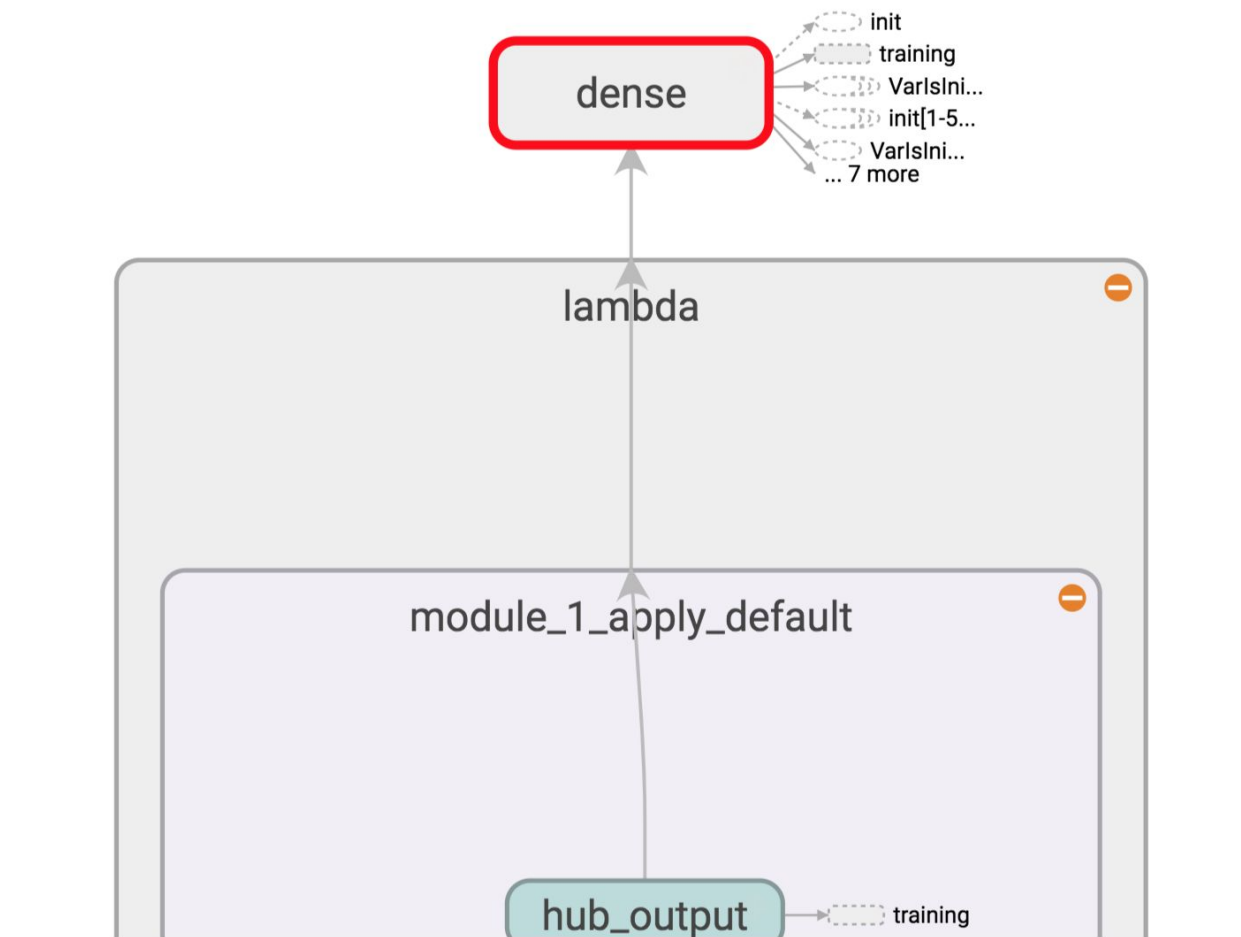
Accuracy: 0.801

TensorBoard Graph:

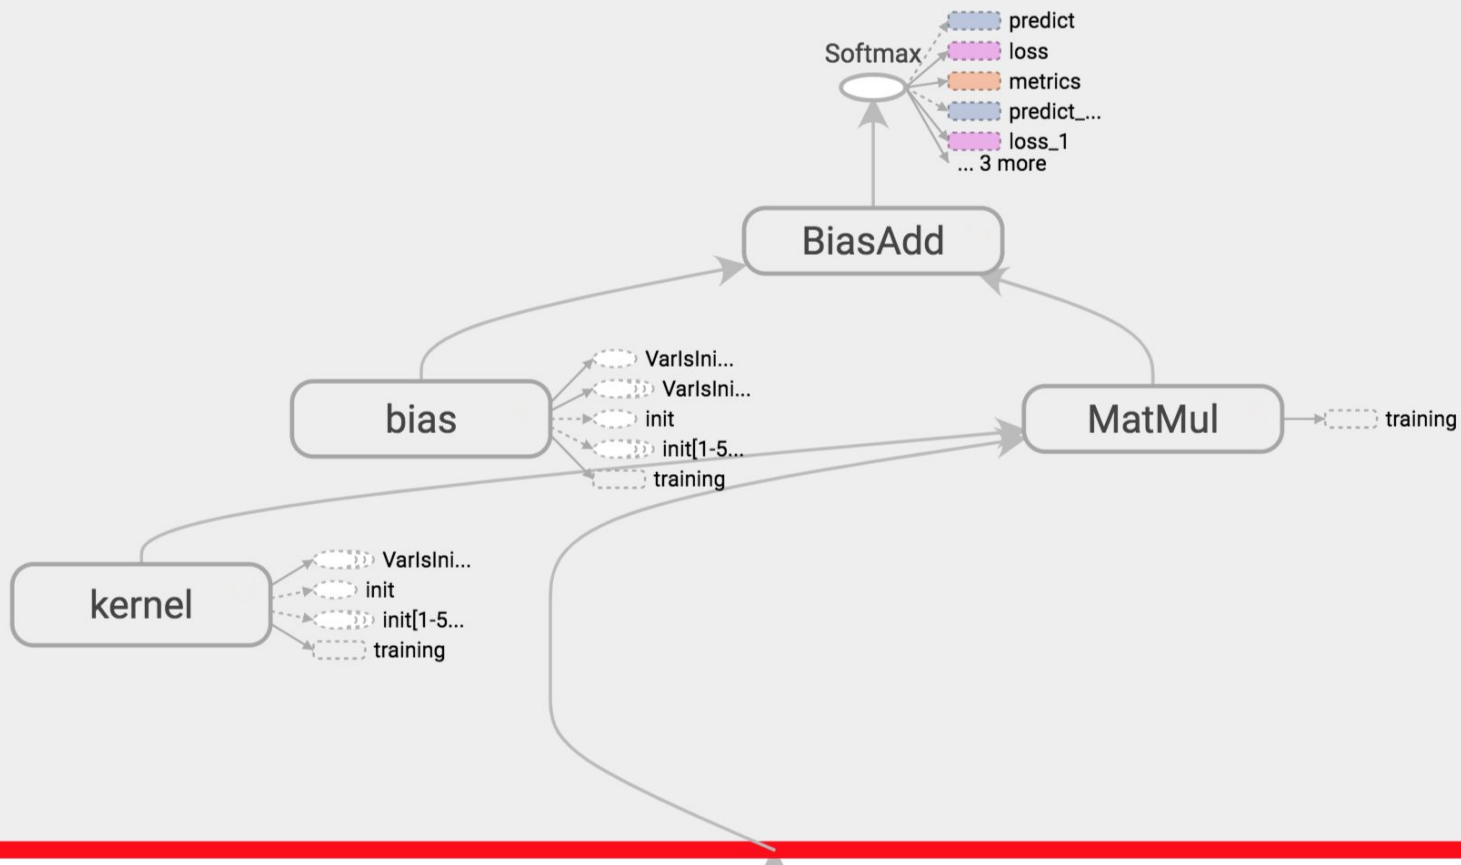








dense



Adding one more dense layer:



```
model = tf.keras.Sequential([
    features_extractor_layer,
    layers.Dense(100, activation='relu'),
    layers.Dense(training_data.num_classes, activation='softmax')
])
model.summary()
```

Layer (type)	Output Shape	Param #
=====		
lambda (Lambda)	(None, 1280)	0
dense (Dense)	(None, 100)	128100
dense_1 (Dense)	(None, 2)	202
=====		

Total params: 128,302

Trainable params: 128,302

Non-trainable params: 0



actual Istanbul Taipei

predicted

	Istanbul	Taipei
Istanbul	436	73
Taipei	72	435

Accuracy: 0.857

Adding one more dense layer:

```
model = tf.keras.Sequential([
    features_extractor_layer,
    layers.Dense(150, activation='relu'),
    layers.Dense(100, activation='relu'),
    layers.Dense(training_data.num_classes, activation='softmax')
])
model.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
lambda (Lambda)	(None, 1280)	0
dense (Dense)	(None, 150)	192150
dense_1 (Dense)	(None, 100)	15100
dense_2 (Dense)	(None, 2)	202
=====	=====	=====

Total params: 207,452

Trainable params: 207,452

Non-trainable params: 0

	actual	Istanbul	Taipei
predicted			
Istanbul		436	73
Taipei		72	435

Accuracy: 0.859